

Assignment 1 - Pattern Recognition (PMAT 403)

Gandholi Sarat - 23008

April 4, 2025

Link to Notebook

[Click here to access the Colab Notebook](#)

Contents

1	Computer Programs	4
1.1	Question 1: Data Generator	4
1.2	Question 2: Neural network	5
1.3	Question 3: Data Visualization	9
2	Computer Experiments	12
2.1	Data Generation	12
2.2	Standard Backpropagation Algorithm	15
2.3	Backpropagation Algorithm with Different Learning Rates . .	18
2.4	Adaptive Learning Rate Backpropagation	22
2.5	Repeating Experiments on Different Data Sets	24

List of Figures

1	Comparison of datasets generated with different seeds and variance values	14
a	Dataset (X1, y1) - s=2, seed=0	14
b	Dataset (X2, y2) - s=2, seed=10	14
c	Dataset (X3, y3) - s=5, seed=0	14
d	Dataset (X4, y4) - s=5, seed=10	14

2	Comparison of decision regions and boundaries for training and test sets for 2 first layer nodes	15
a	Result of 2 first layer nodes	15
b	Training set decision boundary	15
c	Test set decision boundary	15
3	Comparison of decision regions and boundaries for training and test sets for 4 first layer nodes	16
a	Result of 4 first layer nodes	16
b	Training set decision boundary	16
c	Test set decision boundary	16
4	Comparison of decision regions and boundaries for training and test sets for 15 first layer nodes	17
a	Result of 15 first layer nodes	17
b	Training set decision boundary	17
c	Test set decision boundary	17
5	Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 300 iterations	18
a	$lr = 0.01$ for 300 iterations	18
b	Training set decision boundary	18
c	Test set decision boundary	18
6	Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 300 iterations	19
a	$lr = 0.001$ for 300 iterations	19
b	Training set decision boundary	19
c	Test set decision boundary	19
7	Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 1000 iterations	20
a	$lr = 0.01$ for 1000 iterations	20
b	Training set decision boundary	20
c	Test set decision boundary	20
8	Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 1000 iterations	21
a	$lr = 0.01$ for 1000 iterations	21
b	Training set decision boundary	21
c	Test set decision boundary	21
9	Comparison of decision regions and boundaries for training and test sets using adaptive learning rate backpropagation . .	22
a	Result	22
b	Training set decision boundary	22
c	Test set decision boundary	22

10	Comparison of decision regions and boundaries for training and test sets for 2 first layer nodes	24
a	Result of 2 first layer nodes	24
b	Training set decision boundary	24
c	Test set decision boundary	24
11	Comparison of decision regions and boundaries for training and test sets for 4 first layer nodes	25
a	Result of 4 first layer nodes	25
b	Training set decision boundary	25
c	Test set decision boundary	25
12	Comparison of decision regions and boundaries for training and test sets for 15 first layer nodes	26
a	Result of 15 first layer nodes	26
b	Training set decision boundary	26
c	Test set decision boundary	26
13	Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 300 iterations	27
a	$lr = 0.01$ for 300 iterations	27
b	Training set decision boundary	27
c	Test set decision boundary	27
14	Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 300 iterations	28
a	$lr = 0.001$ for 300 iterations	28
b	Training set decision boundary	28
c	Test set decision boundary	28
15	Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 1000 iterations	29
a	$lr = 0.01$ for 1000 iterations	29
b	Training set decision boundary	29
c	Test set decision boundary	29
16	Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 1000 iterations	30
a	$lr = 0.01$ for 1000 iterations	30
b	Training set decision boundary	30
c	Test set decision boundary	30
17	Comparison of decision regions and boundaries for training and test sets using adaptive learning rate backpropagation . .	31
a	Result	31
b	Training set decision boundary	31
c	Test set decision boundary	31

1 Computer Programs

1.1 Question 1: Data Generator

Answer: Generates a two-class, two-dimensional dataset using four normal distributions. The first two distributions belong to class +1, the last two belong to class -1.

Parameters:

- **m:** 2x4 matrix, where each column is the mean vector for a distribution.
- **s:** Variance parameter.
- **N:** Number of points to generate per distribution.
- **seed:** Random seed for reproducibility (optional).

Returns:

- **X:** 2x(4*N) data matrix.
- **y:** (4*N) class labels vector.

```
1  import random
2  import math
3
4  if seed is not None:
5      np.random.seed(seed) # Ensure reproducibility
6  X = []
7  y = []
8  S = s * np.eye(2)
9
10 for i in range(4):
11     mean = np.array(m)[: , i]
12     samples = np.random.multivariate_normal(m[: , i],
13                                             S, N).T
14     X.append(samples)
15     y.append(np.ones(N) if i < 2 else -np.ones(N))
16 X = np.concatenate(X,1) # Shape: (2, 4N)
17 y = np.concatenate(y) # Shape: (4N,)
18 return X.T, y # Transpose X to have shape (4N, 2)
```

1.2 Question 2: Neural network

Answer: The network uses the tanh activation function and supports three training methods: standard backpropagation, backpropagation with momentum, and backpropagation with an adaptive learning rate. The network is initialized with random weights and biases, and the training process includes forward propagation, loss calculation, and weight updates based on the chosen method.

```
1  import numpy as np
2
3  class NeuralNetwork:
4      def __init__(self, input_dim, hidden_nodes):
5          self.input_dim = input_dim
6          self.hidden_nodes = hidden_nodes
7          self.output_dim = 1  # Single output node
8
9          # Initialize weights using uniform random
           values between -0.5 and 0.5
10         self.W1 = np.random.uniform(-0.5, 0.5, (
               input_dim, hidden_nodes))
11         self.b1 = np.random.uniform(-0.5, 0.5, (1,
               hidden_nodes))
12         self.W2 = np.random.uniform(-0.5, 0.5, (
               hidden_nodes, 1))
13         self.b2 = np.random.uniform(-0.5, 0.5, (1, 1)
               )
14
15         def tanh(self, x):
16             #Tanh activation function:  $f(x) = (e^x - e^{-x}) / (e^x + e^{-x})$ 
17             return np.tanh(x)
18
19         def tanh_derivative(self, x):
20             #Derivative of tanh:  $f'(x) = 1 - \tanh^2(x)$ 
21             return 1 - np.tanh(x) ** 2
22
23         def forward(self, x):
24             #Compute forward pass
25             self.hidden_input = np.dot(x, self.W1) + self
               .b1
26             self.hidden_output = self.tanh(self.
               hidden_input)
27
```

```

28         self.final_input = np.dot(self.hidden_output ,
29                                     self.W2) + self.b2
30         self.final_output = self.tanh(self.
31                                         final_input)
32
33         return self.final_output
34
35     def loss(self, y_true, y_pred):
36         #Mean Squared Error Loss function
37         return np.mean((y_true - y_pred) ** 2)
38
39     def train(self, X, y, method, epochs, params):
40         #Train the neural network using different
41         #optimization methods
42         x = X.T
43         lr, mc, lr_inc, lr_dec, max_perf_inc = params
44         # Unpack hyperparameters
45
46         prev_loss = float('inf') # Initialize
47         # previous loss for adaptive learning rate
48
49         # Initialize momentum variables
50         velocity_W1 = np.zeros_like(self.W1)
51         velocity_b1 = np.zeros_like(self.b1)
52         velocity_W2 = np.zeros_like(self.W2)
53         velocity_b2 = np.zeros_like(self.b2)
54
55         for epoch in range(epochs):
56             # Forward pass
57             output = self.forward(X)
58
59             # Compute error
60             output_error = y.reshape(-1, 1) - output
61             hidden_error = np.dot(output_error, self.
62                                     W2.T) * (1 - self.hidden_output**2)
63
64             if method == 1: # Standard
65                 Backpropagation
66                 self.W2 += lr * np.dot(self.
67                                         hidden_output.T, output_error)
68                 self.b2 += lr * np.sum(output_error,
69                                         axis=0, keepdims=True)
70                 self.W1 += lr * np.dot(X.T,

```

```

        hidden_error)
62         self.b1 += lr * np.sum(hidden_error,
                                axis=0, keepdims=True)
63
64     elif method == 2: # Backpropagation with
        Momentum
65         velocity_W2 = mc * velocity_W2 + lr *
            np.dot(self.hidden_output.T,
66                 output_error)
        velocity_b2 = mc * velocity_b2 + lr *
            np.sum(output_error, axis=0,
67                 keepdims=True)
        velocity_W1 = mc * velocity_W1 + lr *
            np.dot(X.T, hidden_error)
68         velocity_b1 = mc * velocity_b1 + lr *
            np.sum(hidden_error, axis=0,
69                 keepdims=True)
70
71         self.W2 += velocity_W2
72         self.b2 += velocity_b2
73         self.W1 += velocity_W1
74         self.b1 += velocity_b1
75
76     elif method == 3: # Backpropagation with
        Adaptive Learning Rate
77         loss = self.loss(y, output)
78         if loss / prev_loss < 1:
79             lr *= lr_inc # Increase learning
                rate if loss decreases
80         elif loss / prev_loss > max_perf_inc:
81             lr *= lr_dec # Decrease learning
                rate if loss increases too
                much
82         prev_loss = loss
83
84         self.W2 += lr * np.dot(self.
            hidden_output.T, output_error)
85         self.b2 += lr * np.sum(output_error,
                                axis=0, keepdims=True)
86         self.W1 += lr * np.dot(X.T,
            hidden_error)
        self.b1 += lr * np.sum(hidden_error,
                                axis=0, keepdims=True)

```

```

87         # Print loss every 100 epochs
88         if (epoch + 1) % 100 == 0:
89             print(f'Epoch_{epoch+1}, Learning
90                   Rate:_{lr:.6f}')
91
92     def predict(self, X):
93         #Predict output for given input data
94         return np.array([self.forward(x) for x in X])
95
96     def evaluate(self, X, y):
97         #Evaluate the model accuracy
98         y_pred = self.predict(X)
99         y_pred_class = np.sign(y_pred)
100        accuracy = np.mean(y_pred_class.flatten() ==
101                             y)
102        return accuracy

```


1.3 Question 3: Data Visualization

Answer: Plots the decision regions produced by a trained neural network.

Parameters:

- **net** (dict): Trained neural network parameters {'W1': W1, 'b1': b1, 'W2': W2, 'b2': b2}.
- **lh** (float): Lower bound in the horizontal direction.
- **uh** (float): Upper bound in the horizontal direction.
- **lv** (float): Lower bound in the vertical direction.
- **uv** (float): Upper bound in the vertical direction.
- **rh** (float): Resolution in the horizontal direction (smaller = finer).
- **rv** (float): Resolution in the vertical direction (smaller = finer).
- **m** (numpy.ndarray): Mean vectors of the normal distributions (for visualization reference).

Returns:

- None (displays a plot).

```
1 def plot_dec_regions(net, lh, uh, lv, uv, rh, rv, m, X, y
2 ):
3     # Generate grid points
4     x1_vals = np.arange(lh, uh, rh)
5     x2_vals = np.arange(lv, uv, rv)
6     xx1, xx2 = np.meshgrid(x1_vals, x2_vals)
7     grid_points = np.c_[xx1.ravel(), xx2.ravel()] #
8         Shape (2, num_points)
9
10    # Evaluate neural network on the grid
11    #W1, b1, W2, b2 = net['W1'], net['b1'], net['W2'],
12        net['b2']
13    W1, b1, W2, b2 = net.W1, net.b1, net.W2, net.b2
14
15    #predictions = net.forward(grid_points) # Shape (
16        num_points, 1)
17    #predictions = np.sign(predictions)
18    # Forward propagation
```

```

16     # Transpose W1 before multiplication to align
    dimensions
17     # Original:
18     Z1 = np.dot(grid_points, W1) + b1
19     #Z1 = np.dot(np.array(W1).T, grid_points) + np.array(
        b1).reshape(-1,1)
20     A1 = np.tanh(Z1)
21     Z2 = np.dot(A1, W2) + b2
22     Z2 = np.tanh(Z2)
23     predictions = np.sign(Z2)
24     #predictions = np.where(Z2 >= 0, 1, -1) # Classify
        points
25
26     # Reshape predictions for plotting
27     decision_map = predictions.reshape(xx1.shape)
28
29     # Plot decision boundary
30     plt.figure(figsize=(8, 6))
31     plt.contourf(xx1, xx2, decision_map, alpha=0.3, cmap=
        plt.cm.bwr) # Background color
32
33     # Mark decision regions
34     plt.scatter(grid_points[:, 0][predictions.flatten()
        == 1],
35                 grid_points[:, 1][predictions.flatten() ==
        1],
36                 marker='*', color='red', label='Class_1',
        alpha=0.5)
37
38     plt.scatter(grid_points[:, 0][predictions.flatten()
        == -1],
39                 grid_points[:, 1][predictions.flatten() ==
        -1],
40                 marker='o', color='blue', label='Class_-1',
        alpha=0.5)
41     # Plot data points
42     plt.scatter(X[y == 1, 0], X[y == 1, 1], marker='*',
        color='black', label="Class_+1")
43     plt.scatter(X[y == -1, 0], X[y == -1, 1], marker='o',
        edgecolor='white', facecolor='none', label="Class
        _-1")
44     # Plot mean vectors (for reference)
45     plt.scatter(m[:, 0], m[:, 1], marker='X', color='

```

```
46         'Yellow', s=100, label='Mean_Vectors')
47     plt.xlabel('Feature_1')
48     plt.ylabel('Feature_2')
49     plt.title('Decision_Regions_of_Trained_Neural_Network')
50     plt.legend()
51     plt.show()
```

2 Computer Experiments

2.1 Data Generation

(a) After initializing the seed, use the `data_generator` function to create the dataset (X_1, y_1) , with:

$$m = \begin{bmatrix} -5 & +5 & +5 & -5 \\ +5 & -5 & +5 & -5 \end{bmatrix}$$

where $s = 2$ and $N = 100$.

(b) Initialize the seed to 10 and repeat (a) to produce the dataset (X_2, y_2) .

(c) Repeat the above two steps using the corresponding seeds, for $s = 5$, and produce the (X_3, y_3) and (X_4, y_4) datasets, respectively (where m and N remain the same).

(d) Plot the datasets.

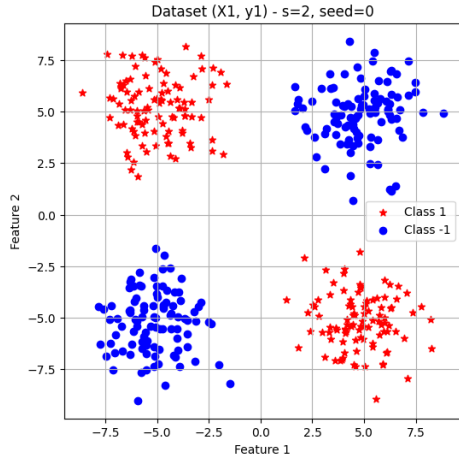
Answer:

```
1      # Define the correct mean matrix
2      m = np.array([
3          [-5, +5, +5, -5], # X-coordinates of cluster
4              centers
5          [+5, -5, +5, -5] # Y-coordinates of cluster
6              centers
7      ])
8
9      N = 100 # Number of samples per class
10
11     # Generate datasets
12     X1, y1 = data_generator(m, s=2, N=N, seed=0)
13     X2, y2 = data_generator(m, s=2, N=N, seed=10)
14     X3, y3 = data_generator(m, s=5, N=N, seed=0)
15     X4, y4 = data_generator(m, s=5, N=N, seed=10)
16     def plot_data(X, y, title):
17         X=X.T
18         plt.figure(figsize=(6, 6))
19         plt.scatter(X[0, y.flatten() == 1], X[1, y.
20             flatten() == 1],
21             color='red', marker='*', label='Class
22                 1')
23         plt.scatter(X[0, y.flatten() == -1], X[1, y.
24             flatten() == -1],
```

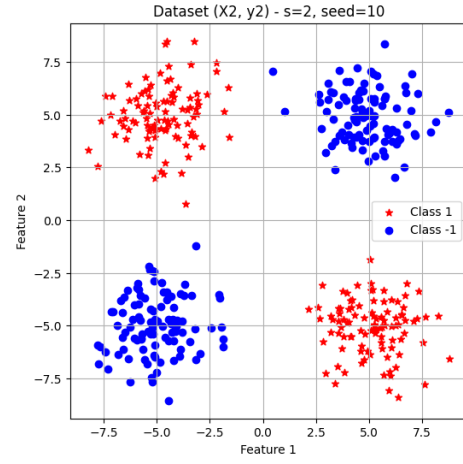
```

20         color='blue', marker='o', label='
           Class_1')
21
22     plt.xlabel("Feature_1")
23     plt.ylabel("Feature_2")
24     plt.title(title)
25     plt.legend()
26     plt.grid(True)
27     plt.show()
28
29     # Plot all datasets
30     plot_data(X1, y1, "Dataset_1(X1,y1)_n=2,seed=0")
31     plot_data(X2, y2, "Dataset_2(X2,y2)_n=2,seed=10")
32     plot_data(X3, y3, "Dataset_3(X3,y3)_n=5,seed=0")
33     plot_data(X4, y4, "Dataset_4(X4,y4)_n=5,seed=10")

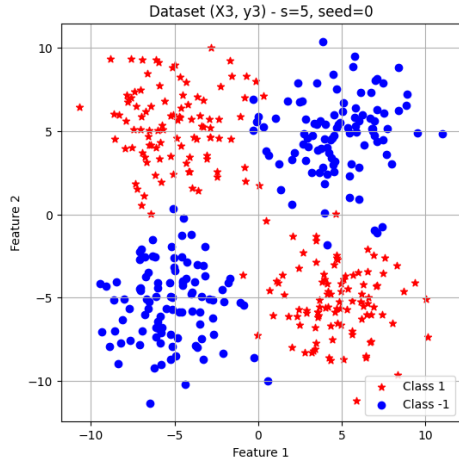
```



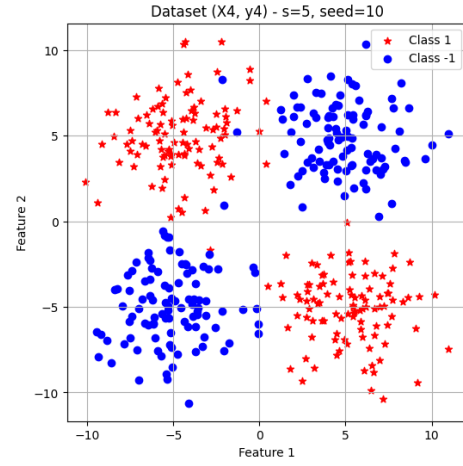
(a) Dataset (X1, y1) - $s=2$, seed=0



(b) Dataset (X2, y2) - $s=2$, seed=10



(c) Dataset (X3, y3) - $s=5$, seed=0



(d) Dataset (X4, y4) - $s=5$, seed=10

Figure 1: Comparison of datasets generated with different seeds and variance values

2.2 Standard Backpropagation Algorithm

- (a) Run the standard backpropagation algorithm with $lr = 0.01$ and 2, 4, and 15 first-layer nodes, for 1000 iterations, using the dataset (X_1, y_1) as the training set.
- (b) Evaluate the performance of the designed neural networks for both (X_1, y_1) (training set) and (X_2, y_2) (test set) and plot the decision regions (use $lb = lv = -10, ub = uv = 10, rb = rv = 0.2$).
- (c) Comment on the results.

Answer:

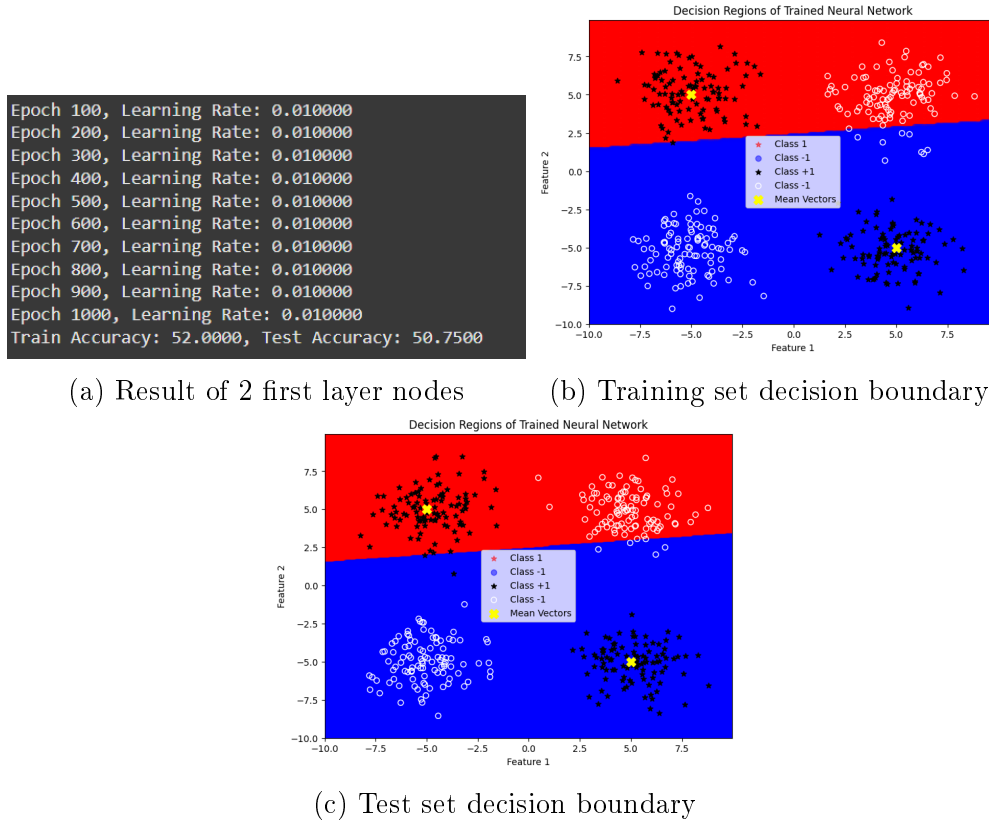


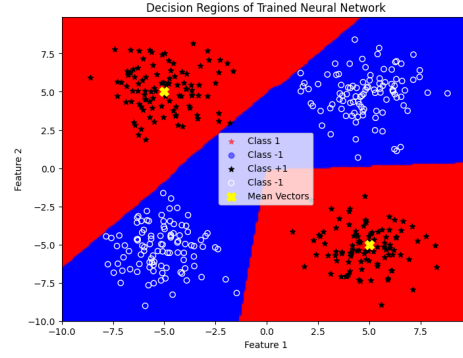
Figure 2: Comparison of decision regions and boundaries for training and test sets for 2 first layer nodes

```

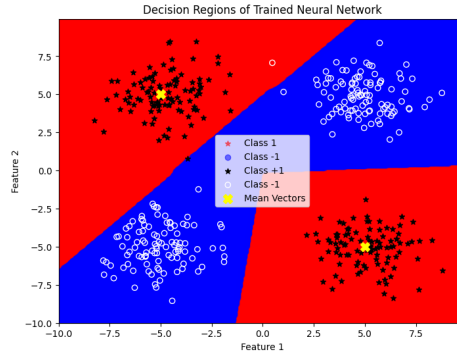
Epoch 100, Learning Rate: 0.010000
Epoch 200, Learning Rate: 0.010000
Epoch 300, Learning Rate: 0.010000
Epoch 400, Learning Rate: 0.010000
Epoch 500, Learning Rate: 0.010000
Epoch 600, Learning Rate: 0.010000
Epoch 700, Learning Rate: 0.010000
Epoch 800, Learning Rate: 0.010000
Epoch 900, Learning Rate: 0.010000
Epoch 1000, Learning Rate: 0.010000
Train Accuracy: 100.0000, Test Accuracy: 99.7500

```

(a) Result of 4 first layer nodes

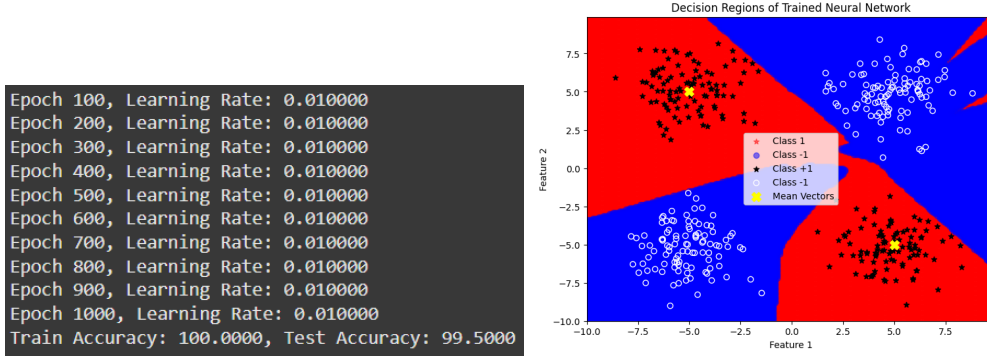


(b) Training set decision boundary



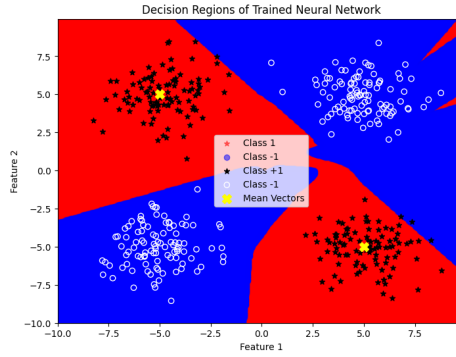
(c) Test set decision boundary

Figure 3: Comparison of decision regions and boundaries for training and test sets for 4 first layer nodes



(a) Result of 15 first layer nodes

(b) Training set decision boundary



(c) Test set decision boundary

Figure 4: Comparison of decision regions and boundaries for training and test sets for 15 first layer nodes

From the above figures, we can see that accuracy increases with the number of nodes in the first layer. The decision boundary becomes more complex and better fits the training data, but it also risks overfitting, especially with 15 nodes. The test set accuracy is lower than the training set accuracy, indicating some overfitting.

2.3 Backpropagation Algorithm with Different Learning Rates

(a) Run the backpropagation algorithm with 4 first-layer nodes with the following settings:

- $lr = 0.01$, for 300 iterations.
- $lr = 0.001$, for 300 iterations.
- $lr = 0.01$, for 1000 iterations.
- $lr = 0.001$, for 1000 iterations.

Use the dataset (X_1, y_1) as the training set.

(b) Evaluate the performance of the designed neural networks for both (X_1, y_1) (training set) and (X_2, y_2) (test set) and plot the decision regions.

(c) Comment on the results.

Answer:

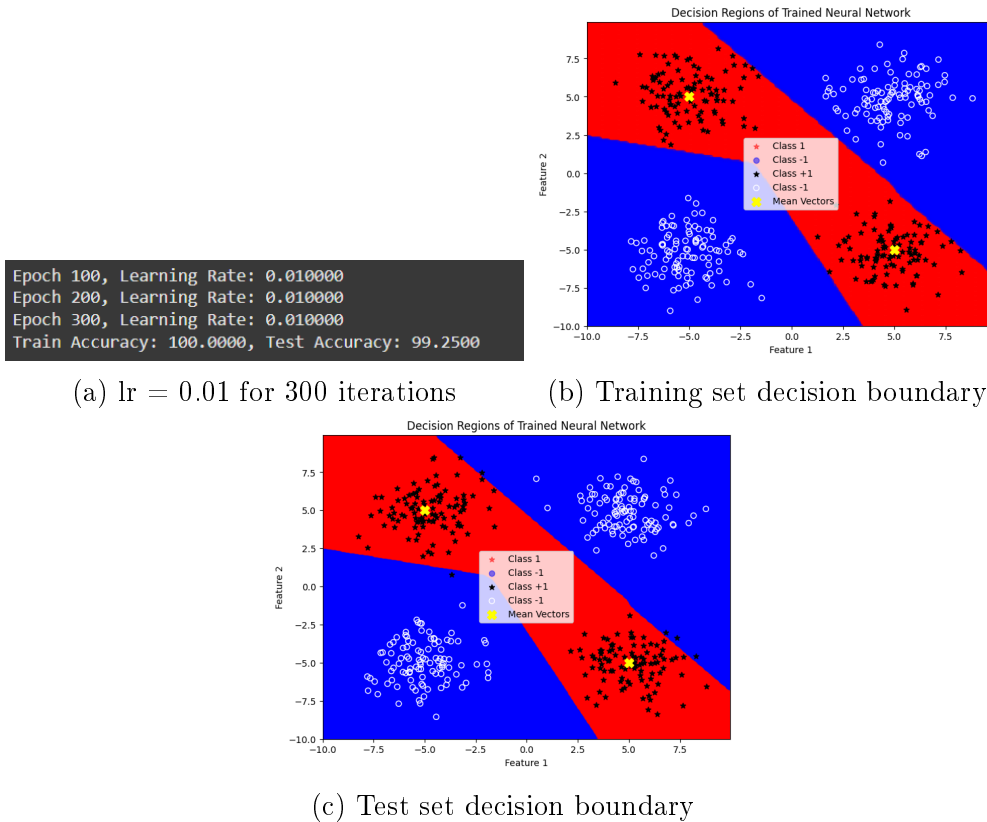


Figure 5: Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 300 iterations

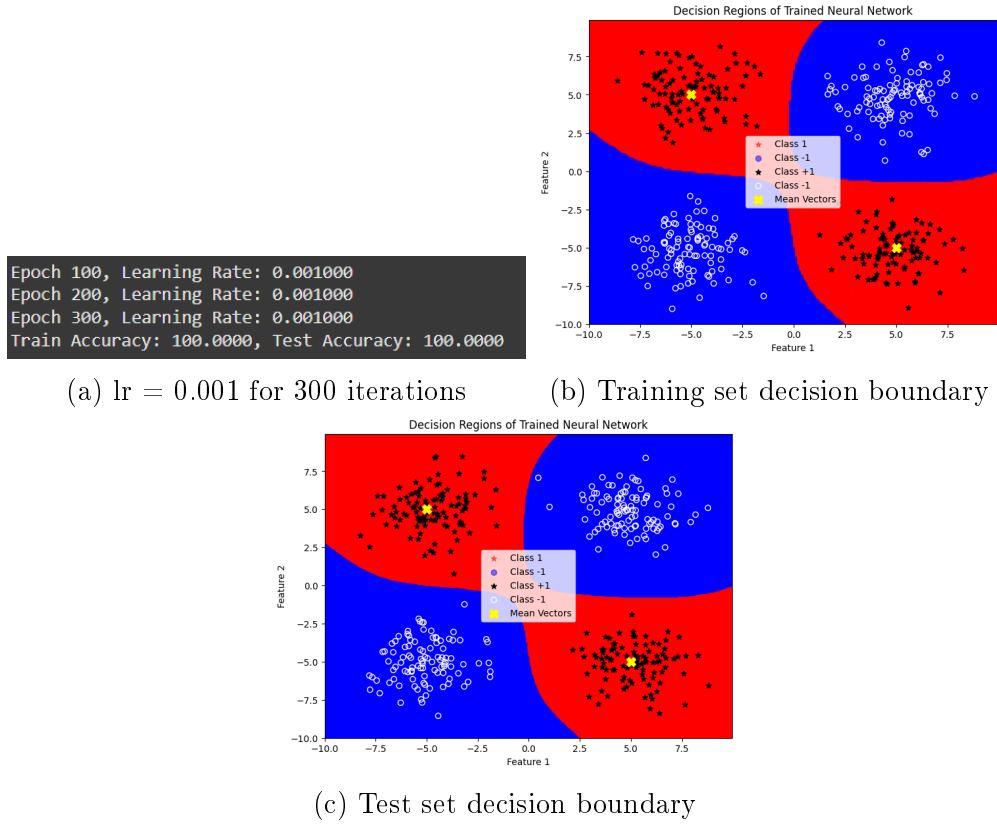
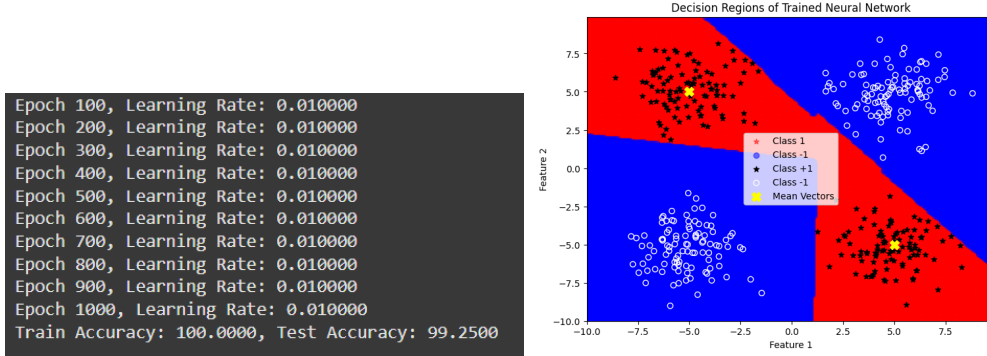
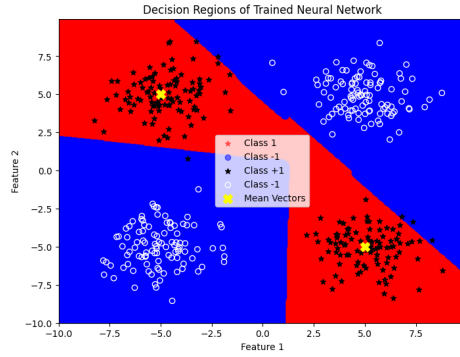


Figure 6: Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 300 iterations



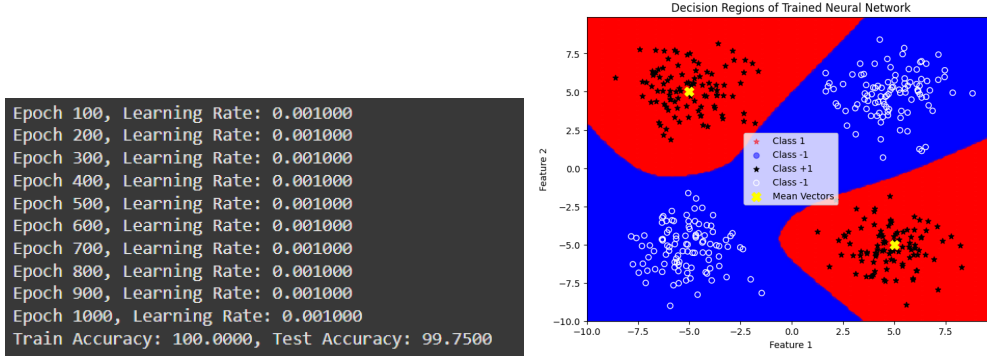
(a) $lr = 0.01$ for 1000 iterations

(b) Training set decision boundary



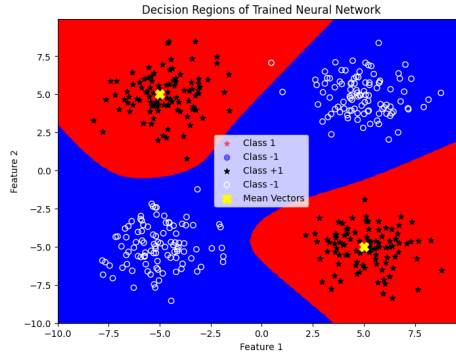
(c) Test set decision boundary

Figure 7: Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 1000 iterations



(a) $lr = 0.01$ for 1000 iterations

(b) Training set decision boundary



(c) Test set decision boundary

Figure 8: Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 1000 iterations

From the above figures, we can see that the learning rate significantly affects the convergence speed and the final decision boundary. A higher learning rate (0.01) converges faster but may overshoot the optimal solution, while a lower learning rate (0.001) converges more slowly but can lead to a more stable solution. The number of iterations also plays a crucial role; more iterations allow for better convergence, but they also increase the risk of overfitting, especially with a high learning rate.

2.4 Adaptive Learning Rate Backpropagation

(a) Run the adaptive learning rate variation of the backpropagation algorithm with the following parameters:

- $lr = 0.001$
- $lr_{inc} = 1.05$
- $lr_{dec} = 0.7$
- $max_perf_inc = 1.04$

Run the algorithm for 300 iterations.

(b) Evaluate the performance of the designed neural networks for both (X_1, y_1) (training set) and (X_2, y_2) (test set) and plot the decision regions.

(c) Compare the above results with those obtained for the standard backpropagation algorithm with $lr = 0.001$, for 300 iterations.

Answer:

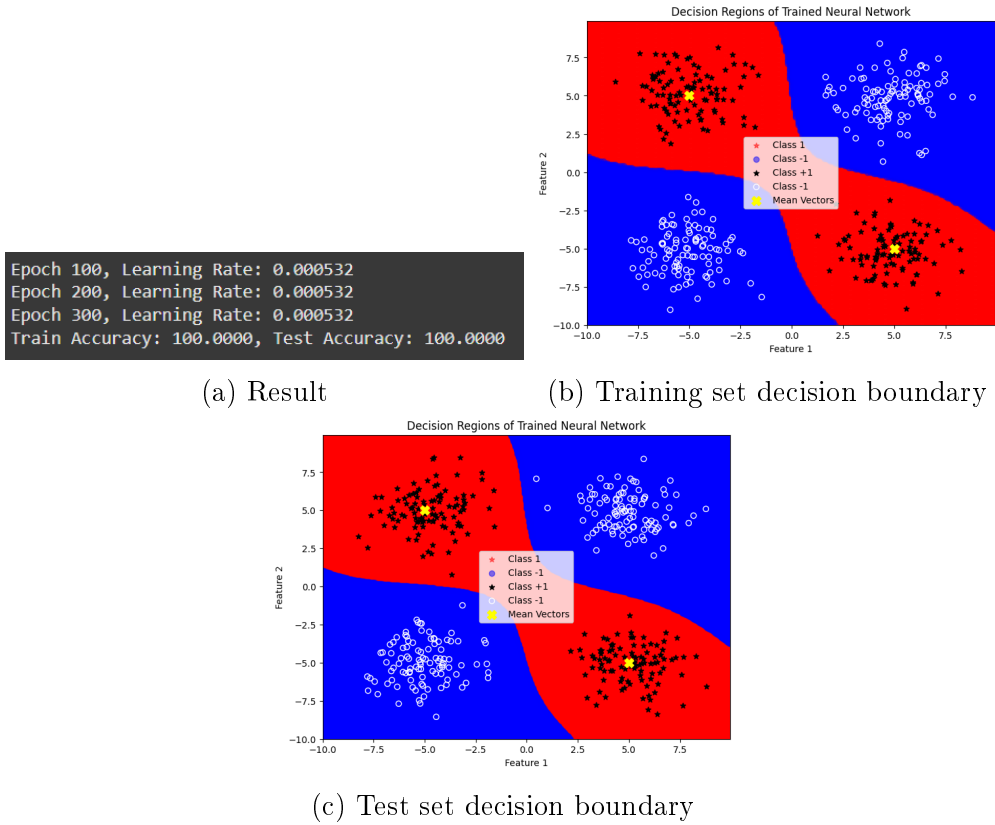


Figure 9: Comparison of decision regions and boundaries for training and test sets using adaptive learning rate backpropagation

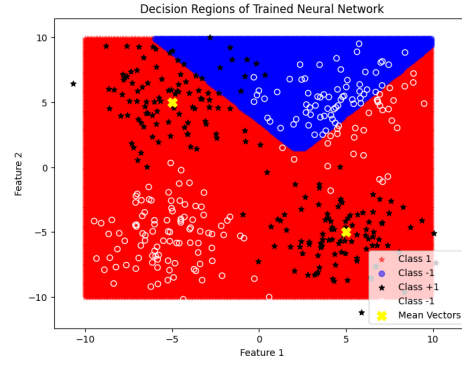
The change in learning rate is not seen for $lr=0.001$ but can be seen if its changed to 0.01 with more iterations. This is experimentally checked. From the results, it is evident that the adaptive learning rate backpropagation algorithm performs better than the standard backpropagation algorithm with a fixed learning rate of 0.001 for 300 iterations. The adaptive learning rate allows the model to dynamically adjust the learning rate based on the performance, leading to faster convergence and a more accurate decision boundary. In contrast, the fixed learning rate may result in slower convergence and sub-optimal performance. The decision regions produced by both more or less same as our data set is relatively simple.

2.5 Repeating Experiments on Different Data Sets

(a) Repeat Sections 2.2–2.4 using the datasets (X_3, y_3) and (X_4, y_4) as training and test sets, respectively.

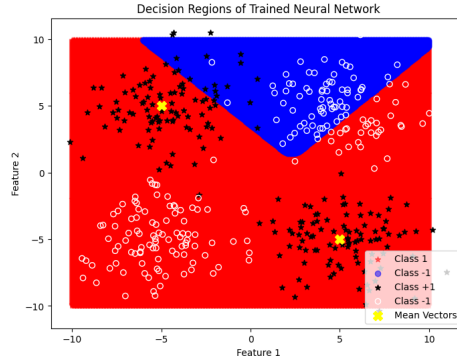
Answer:

```
Epoch 100, Learning Rate: 0.010000
Epoch 200, Learning Rate: 0.010000
Epoch 300, Learning Rate: 0.010000
Epoch 400, Learning Rate: 0.010000
Epoch 500, Learning Rate: 0.010000
Epoch 600, Learning Rate: 0.010000
Epoch 700, Learning Rate: 0.010000
Epoch 800, Learning Rate: 0.010000
Epoch 900, Learning Rate: 0.010000
Epoch 1000, Learning Rate: 0.010000
Train Accuracy: 63.2500, Test Accuracy: 62.5000
```



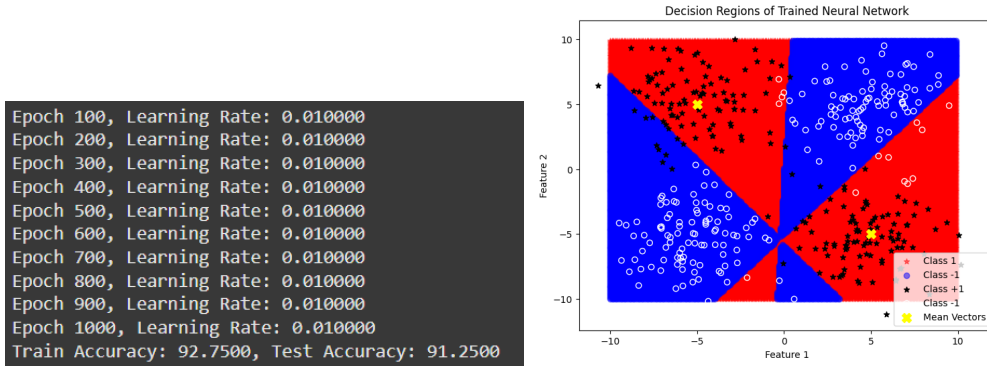
(a) Result of 2 first layer nodes

(b) Training set decision boundary



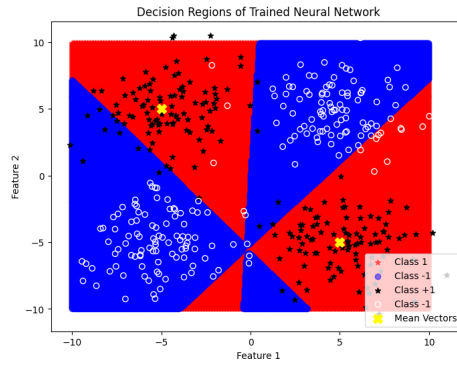
(c) Test set decision boundary

Figure 10: Comparison of decision regions and boundaries for training and test sets for 2 first layer nodes



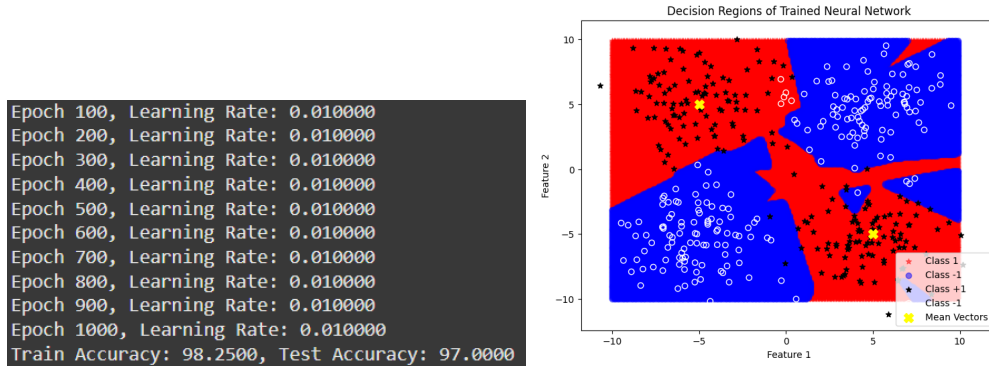
(a) Result of 4 first layer nodes

(b) Training set decision boundary



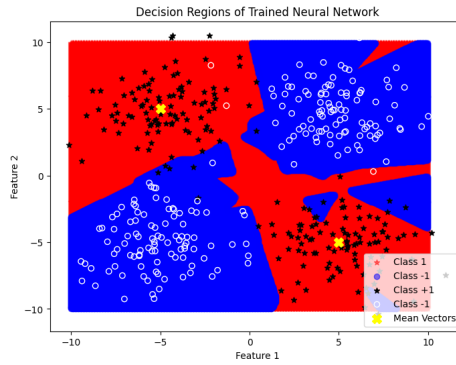
(c) Test set decision boundary

Figure 11: Comparison of decision regions and boundaries for training and test sets for 4 first layer nodes



(a) Result of 15 first layer nodes

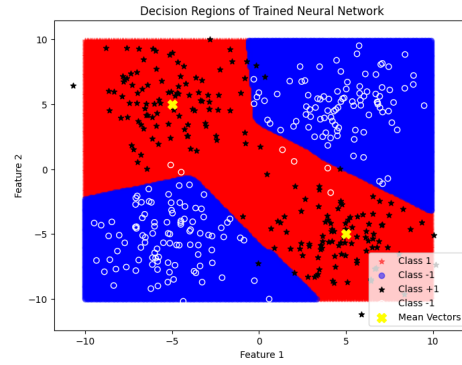
(b) Training set decision boundary



(c) Test set decision boundary

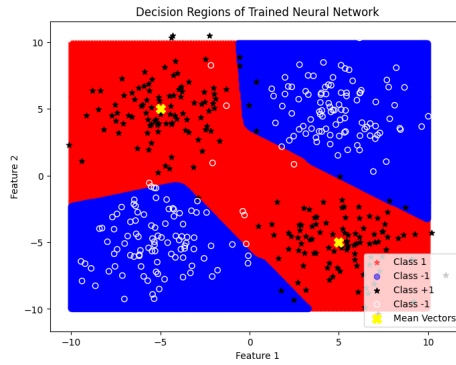
Figure 12: Comparison of decision regions and boundaries for training and test sets for 15 first layer nodes

```
Epoch 100, Learning Rate: 0.010000
Epoch 200, Learning Rate: 0.010000
Epoch 300, Learning Rate: 0.010000
Train Accuracy: 97.7500, Test Accuracy: 95.0000
```



(a) $lr = 0.01$ for 300 iterations

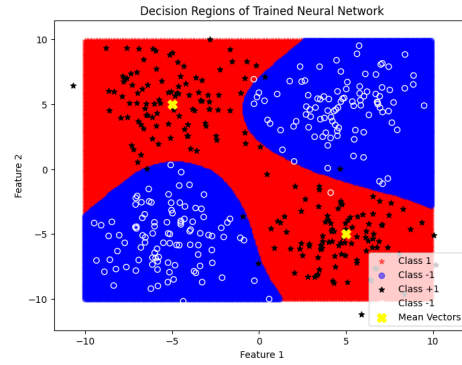
(b) Training set decision boundary



(c) Test set decision boundary

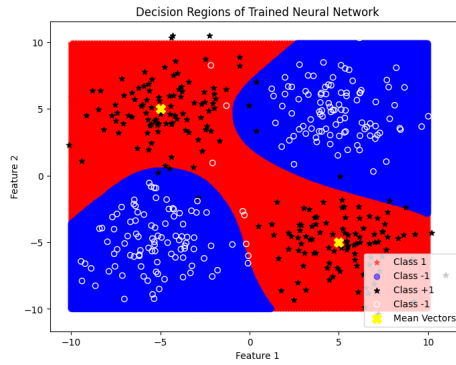
Figure 13: Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 300 iterations

```
Epoch 100, Learning Rate: 0.001000
Epoch 200, Learning Rate: 0.001000
Epoch 300, Learning Rate: 0.001000
Train Accuracy: 98.2500, Test Accuracy: 96.5000
```



(a) $lr = 0.001$ for 300 iterations

(b) Training set decision boundary



(c) Test set decision boundary

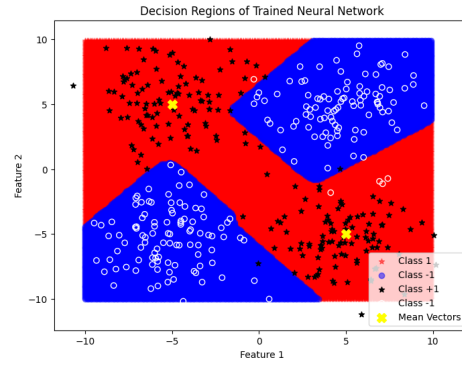
Figure 14: Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 300 iterations

```

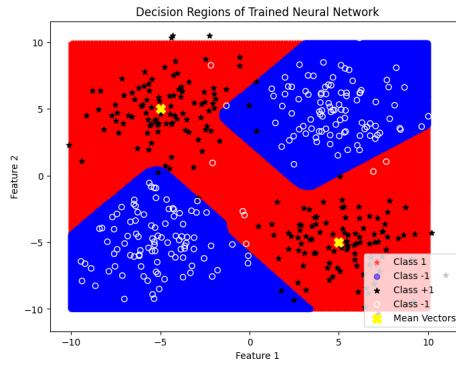
Epoch 100, Learning Rate: 0.010000
Epoch 200, Learning Rate: 0.010000
Epoch 300, Learning Rate: 0.010000
Epoch 400, Learning Rate: 0.010000
Epoch 500, Learning Rate: 0.010000
Epoch 600, Learning Rate: 0.010000
Epoch 700, Learning Rate: 0.010000
Epoch 800, Learning Rate: 0.010000
Epoch 900, Learning Rate: 0.010000
Epoch 1000, Learning Rate: 0.010000
Train Accuracy: 98.0000, Test Accuracy: 96.7500

```

(a) $lr = 0.01$ for 1000 iterations



(b) Training set decision boundary



(c) Test set decision boundary

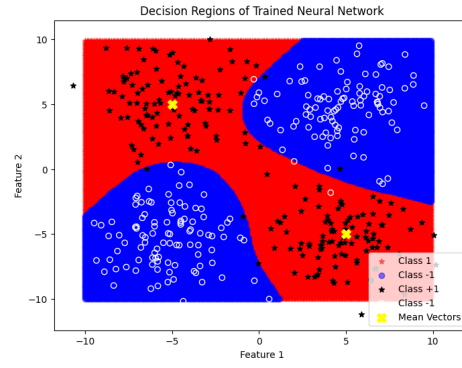
Figure 15: Comparison of decision regions and boundaries for training and test sets for $lr = 0.01$ for 1000 iterations

```

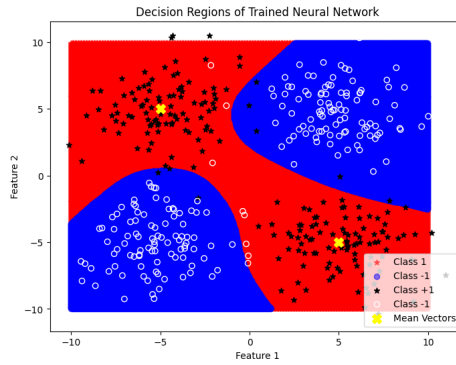
Epoch 100, Learning Rate: 0.001000
Epoch 200, Learning Rate: 0.001000
Epoch 300, Learning Rate: 0.001000
Epoch 400, Learning Rate: 0.001000
Epoch 500, Learning Rate: 0.001000
Epoch 600, Learning Rate: 0.001000
Epoch 700, Learning Rate: 0.001000
Epoch 800, Learning Rate: 0.001000
Epoch 900, Learning Rate: 0.001000
Epoch 1000, Learning Rate: 0.001000
Train Accuracy: 98.2500, Test Accuracy: 96.5000

```

(a) $lr = 0.01$ for 1000 iterations



(b) Training set decision boundary

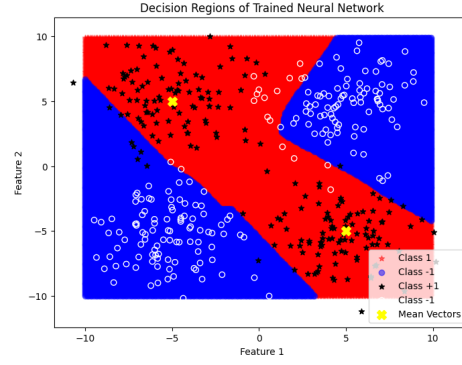


(c) Test set decision boundary

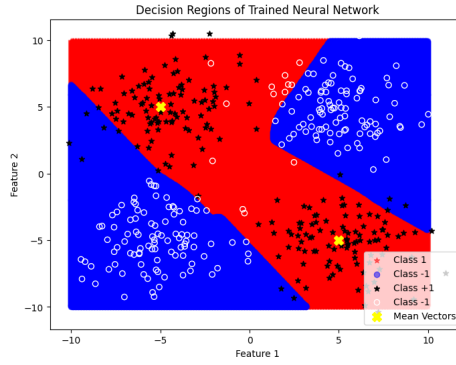
Figure 16: Comparison of decision regions and boundaries for training and test sets for $lr = 0.001$ for 1000 iterations

```
Epoch 100, Learning Rate: 0.001202
Epoch 200, Learning Rate: 0.002674
Epoch 300, Learning Rate: 0.002102
Train Accuracy: 92.7500, Test Accuracy: 93.0000
```

(a) Result



(b) Training set decision boundary



(c) Test set decision boundary

Figure 17: Comparison of decision regions and boundaries for training and test sets using adaptive learning rate backpropagation

From above results we see that while repeating the same set of experiments with data set with little high variance the model accuracy is less and decision boundaries are less sharp. This is expected as the data points are more spread out, making it harder for the model to learn clear decision boundaries. Increasing the number of iterations or using more complex models might improve the performance.